

EN3160 Assignment 2: Fitting and Alignment

Name: Ranaweera R.P.D.D.H.

Index: 220511N

October 14, 2025

Project Repository: Fitting-and-Alignment

Question 1: Blob Detection on Sunflower Field

```
sigma_values = np.linspace(sigma_min, sigma_max, num_sigma)
log_images = []
for sigma in sigma_values:
    size = int(2*np.ceil(4*sigma)+1)
    blur = cv.GaussianBlur(gray, (size, size), sigma)
    log = cv.Laplacian(blur, cv.CV_64F, ksize=1)
    log_images.append(sigma**2*np.abs(log))
log_stack = np.stack(log_images, axis=-1)
local_max = log_stack == maximum_filter(log_stack, size=(3,3,2))
blob_mask = local_max & (log_stack > 0.05)
blobs = [(x,y,i,sigma_values[i],log_stack[y,x,i])
          for y,x,i in np.argwhere(blob_mask)]
blobs_sorted = sorted(blobs, key=lambda b:b[4], reverse=True)
output = im.copy()
for x,y,i,sigma,_ in blobs_sorted[:min(100,len(blobs_sorted))]:
    cv.circle(output, (int(x),int(y)), int(np.sqrt(2)*sigma),
              (0,0,255), 2)
```

Figure 1: Blob Detection using Scale-space LoG

Detected Sunflower Centers using LoG Blob Detection



Figure 2: Detected Blobs over Multiple Scales

LoG Scale-Space Parameters

σ range	[2, 20]
Scales	50
Shape	(360, 360, 50)
Total blobs	4611
Threshold	0.05
Neighborhood	(3, 3, 2)

Top 5 Largest Circles (by radius)

Rank	X	Y	σ	Radius	Response
1	47	0	20.000	28.28	0.102926
2	295	0	20.000	28.28	0.069862
3	68	31	20.000	28.28	0.075601
4	69	33	20.000	28.28	0.075608
5	310	50	20.000	28.28	0.065867

Question 2: Line and Circle Fitting with RANSAC

Line Estimation using RANSAC

a	0.71
b	0.70
d	-1.48
Number of inliers	35

Circle Estimation using RANSAC

Center (x_0, y_0)	(2.12, 2.35)
Radius r	9.88
Number of inliers	52

(d) Circle-First Approach Analysis

If we fit the circle first, many line points appear as outliers and distort the estimate. The circle fitting RANSAC would try to include them, causing incorrect center and radius, poor inlier classification, and degraded line fitting due to misclassified points. Thus, it's better to fit the line first, remove its inliers, then fit the circle.

Listing 1: RANSAC Line Fitting

```
def fit_line(points):
    """Fit line  $ax + by + d = 0$ 
    through two points."""
    (x1, y1), (x2, y2) = points
    a = y1 - y2
    b = x2 - x1
    d = -(a * x1 + b * y1)
    norm = np.sqrt(a**2 + b**2)
    return a/norm, b/norm, d/norm

def line_distance(a, b, d, X):
    """Compute normal distance from
    line to all points."""
    return np.abs(a*X[:,0] + b*X[:,1]
    + d)

def ransac_line(X, n_iter=1000,
    threshold=0.6, min_inliers=35):
    best_inliers = []
    best_model = None
    for _ in range(n_iter):
        sample = X[np.random.choice(len(
            X), 2, replace=False)]
        a, b, d = fit_line(sample)
        distances = line_distance(a, b,
            d, X)
        inliers = X[distances <
            threshold]
        if len(inliers) > len(
            best_inliers):
            best_inliers = inliers
            best_model = (a, b, d)
    return best_model, best_inliers
```

Listing 2: RANSAC Circle Fitting

```
def fit_circle(pts):
    """Fit a circle through 3 points."""
    A = np.array([[2*(pts[1,0]-pts[0,0]), 2*(pts[1,1]-
        pts[0,1])],
        [2*(pts[2,0]-pts[0,0]), 2*(pts[2,1]-
        pts[0,1])]])
    b = np.array([(pts[1,0]**2 - pts[0,0]**2) +
        (pts[1,1]**2 - pts[0,1]**2),
        (pts[2,0]**2 - pts[0,0]**2) +
        (pts[2,1]**2 - pts[0,1]**2)])
    center = np.linalg.solve(A, b)
    radius = np.sqrt((pts[0,0]-center[0])**2 +
        (pts[0,1]-center[1])**2)
    return center, radius

def ransac_circle(X, n_iter=1000, threshold=1.5,
    min_inliers=45):
    best_inliers = []
    best_model = None
    for _ in range(n_iter):
        sample = X[np.random.choice(len(X), 3, replace
            =False)]
        try:
            center, r = fit_circle(sample)
        except np.linalg.LinAlgError:
            continue
        dist = np.sqrt((X[:,0]-center[0])**2 + (X
           [:,1]-center[1])**2)
        inliers = X[np.abs(dist - r) < threshold]
        if len(inliers) > len(best_inliers):
            best_inliers = inliers
            best_model = (center, r)
    return best_model, best_inliers
```

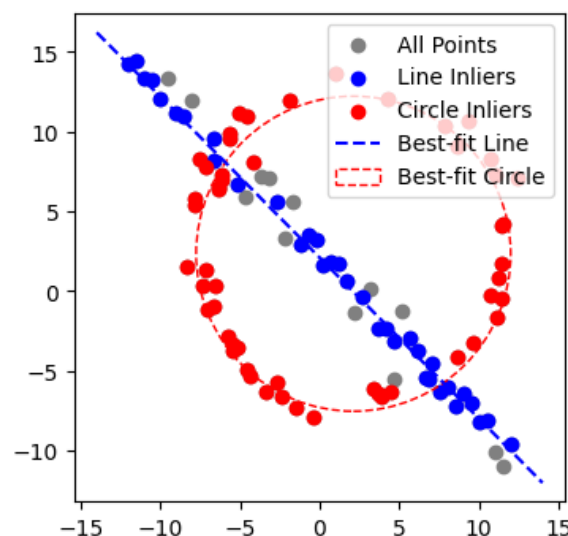


Figure 3: RANSAC line and circle fitting with inliers and sample points

Question 3: Homography and Flag Superimposition

Listing 3: Homography Computation and Warping

```

1 def compute_homography(src_pts, dst_pts):
2     assert src_pts.shape[0] == dst_pts.shape[0] >=
        4
3     n = src_pts.shape[0]
4     A = []
5     for i in range(n):
6         x, y = src_pts[i]
7         u, v = dst_pts[i]
8         A.append([-x, -y, -1, 0, 0, 0, u*x, u*y, u
                ])
9         A.append([0, 0, 0, -x, -y, -1, v*x, v*y, v
                ])
10    A = np.array(A)
11    U, S, Vt = np.linalg.svd(A)
12    H = Vt[-1].reshape(3, 3)
13    H = H / H[2, 2]
14    return H
15
16 def warp_image(src_img, H, dst_shape):
17     h, w = dst_shape[:2]
18     warped = np.zeros((h, w, 3), dtype=np.uint8)
19     mask = np.zeros((h, w), dtype=np.uint8)
20     H_inv = np.linalg.inv(H)
21     y_coords, x_coords = np.mgrid[0:h, 0:w]
22     ones = np.ones_like(x_coords)
23     dst_coords = np.stack([x_coords, y_coords, ones
        ], axis=-1)
24     dst_coords_flat = dst_coords.reshape(-1, 3).T
25     src_coords_flat = H_inv @ dst_coords_flat
26     src_x = src_coords_flat[0] / src_coords_flat[2]
27     src_y = src_coords_flat[1] / src_coords_flat[2]
28     src_x = src_x.reshape(h, w)
29     src_y = src_y.reshape(h, w)
30     src_h, src_w = src_img.shape[:2]
31     valid_mask = (src_x >= 0) & (src_x < src_w - 1)
        & \
32         (src_y >= 0) & (src_y < src_h - 1)
33     x0 = np.floor(src_x).astype(int); x1 = x0 + 1
34     y0 = np.floor(src_y).astype(int); y1 = y0 + 1
35     x0 = np.clip(x0, 0, src_w - 1); x1 = np.clip(x1
        , 0, src_w - 1)

```

Listing 4: Blending and Processing


```

39 y0 = np.clip(y0, 0, src_h - 1); y1 = np.clip(y1
        , 0, src_h - 1)
40 fx = src_x - x0; fy = src_y - y0
41 for c in range(3):
42     wa = (1 - fx) * (1 - fy); wb = fx * (1 - fy
        )
43     wc = (1 - fx) * fy; wd = fx * fy
44     warped[:, :, c] = (wa * src_img[y0, x0, c]
        +
45         wb * src_img[y0, x1, c] +
46         wc * src_img[y1, x0, c] +
47         wd * src_img[y1, x1, c])
48     mask[valid_mask] = 255
49     return warped, mask
50
51 def blend_images(dst_img, src_warped, mask, alpha
    =1.0):
52     result = dst_img.copy()
53     mask_3channel = cv.cvtColor(mask, cv.
        COLOR_GRAY2BGR) / 255.0
54     result = (1 - alpha * mask_3channel) * result +
        alpha * mask_3channel * src_warped
55     return result.astype(np.uint8)
56
57 def process_image_pair(src_img, dst_img,
    dst_points=None, alpha=1.0):
58     h_src, w_src = src_img.shape[:2]
59     src_pts = np.array([[0,0],[w_src-1,0],[w_src-1,
        h_src-1],[0,h_src-1]], dtype=np.float32)
60     dst_pts = np.array(dst_points, dtype=np.float32)
61     if dst_points is not None else
        select_points_interactive(dst_img)
62     H = compute_homography(src_pts, dst_pts)
63     warped, mask = warp_image(src_img, H, dst_img.
        shape)
64     result = blend_images(dst_img, warped, mask,
        alpha)
        return result, H, dst_pts, warped, mask

```



Homography:
$$\begin{bmatrix} 0.7950 & 0.0005649 & 88.0 \\ -0.001277 & 0.4167 & 103.0 \\ -1.2397 \times 10^{-5} & 6.4197 \times 10^{-6} & 1 \end{bmatrix}$$



Homography:
$$\begin{bmatrix} 0.2201 & -0.08589 & 334.0 \\ 0.06482 & 0.11898 & 347.0 \\ 1.3509 \times 10^{-4} & -1.8554 \times 10^{-4} & 1 \end{bmatrix}$$

Question 4: Image Stitching

Listing 5: SIFT Feature Matching and RANSAC

```

1 def compute_sift_matches(img1, img2):
2     sift=cv.SIFT_create()
3     kp1, des1=sift.detectAndCompute(img1, None)
4     kp2, des2=sift.detectAndCompute(img2, None)
5     flann=cv.FlannBasedMatcher(dict(algorithm=1,
6     trees=5), dict(checks=50))
7     matches=flann.knnMatch(des1, des2, k=2)
8     good_matches=[m for m, n in matches if m.
9     distance<0.75*n.distance]
10    print(f"Number of good matches: {len(
11    good_matches)}")
12    return kp1, kp2, good_matches
13
14 def ransac_homography(kp1, kp2, matches, threshold
15 =5.0, iterations=2000):
16     if len(matches)<4: raise ValueError("Not enough
17     matches")
18     pts1=np.float32([kp1[m.queryIdx].pt for m in
19     matches])
20     pts2=np.float32([kp2[m.trainIdx].pt for m in
21     matches])
22     best_H, max_inliers, best_mask=None, 0, None
23     for _ in range(iterations):
24         idx=np.random.choice(len(matches), 4, replace
25         =False)
26         H,_=cv.findHomography(pts1[idx], pts2[idx
27         ], 0)
28         if H is None: continue
29         pts1_h=np.hstack((pts1, np.ones((pts1.shape
30         [0], 1))))
31         pts2_proj=(H@pts1_h.T).T


```

Listing 6: Image Stitching

```

34     pts2_proj=pts2_proj[:, :2]/pts2_proj[:, 2, np.
35     newaxis]
36     inliers=np.linalg.norm(pts2-pts2_proj, axis
37     =1)<threshold
38     if inliers.sum()>max_inliers:
39         max_inliers, best_H, best_mask=inliers.
40         sum(), H, inliers
41     print(f"RANSAC: {max_inliers} inliers found")
42     return best_H, best_mask
43
44 def stitch_images(img1, img2, H):
45     h2, w2=img2.shape[:2]; h1, w1=img1.shape[:2]
46     corners_img1=np.float32([[0, 0], [0, h1], [w1, h1], [
47     w1, 0]]).reshape(-1, 1, 2)
48     warped_corners=cv.perspectiveTransform(
49     corners_img1, H)
50     all_corners=np.vstack((warped_corners, np.
51     float32([[0, 0], [0, h2], [w2, h2], [w2, 0]]).
52     reshape(-1, 1, 2)))
53     xmin, ymin=np.int32(all_corners.min(axis=0).
54     ravel())
55     xmax, ymax=np.int32(all_corners.max(axis=0).
56     ravel())
57     t=[-xmin, -ymin]
58     H_translate=np.array([[1, 0, t[0]], [0, 1, t
59     [1]], [0, 0, 1]])
60     result=cv.warpPerspective(img1, H_translate@H, (
61     xmax-xmin, ymax-ymin))
62     result[t[1]:h2+t[1], t[0]:w2+t[0]]=img2
63     return result

```



Inlier Matches After RANSAC



Stitched Result

Homography matrix used for stitching:

$$\begin{bmatrix} 0.2201 & -0.08589 & 334.0 \\ 0.06482 & 0.11898 & 347.0 \\ 1.3509 \times 10^{-4} & -1.8554 \times 10^{-4} & 1 \end{bmatrix}$$


If the given Homography matrix is used, the following result is obtained:

Intermediate Warped Image (using known H)



Warped Source Image

Final Stitched Result (using known H)



Final Stitched Result